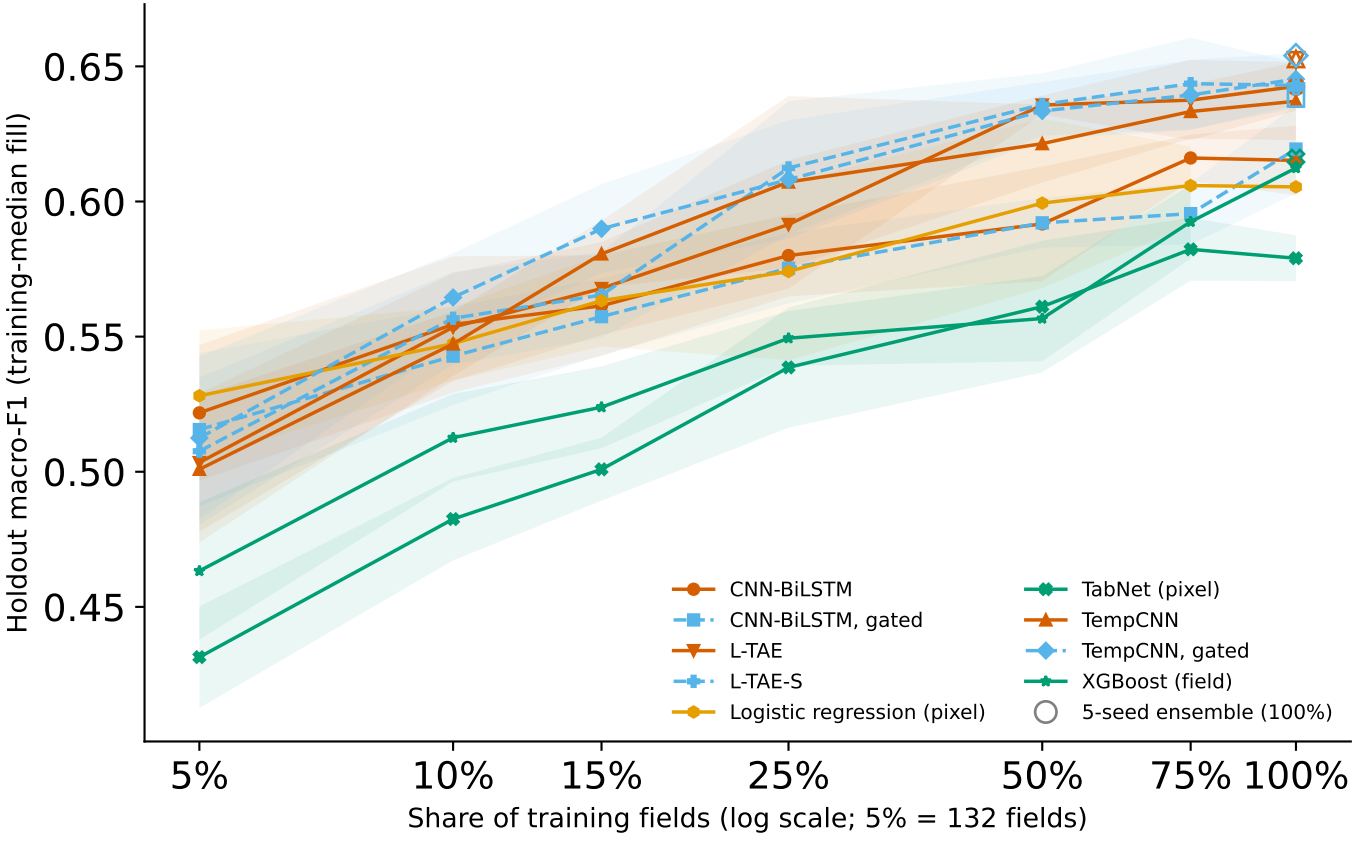
Holdout macro-F1 (training-median fill)
0.65
0.60
0.55
0.50
0.45
5%
10%
15%
25%
50%
75%
100%
Share of training fields (log scale; 5% = 132 fields)
CNN-BiLSTM
CNN-BiLSTM, gated
L-TAE
L-TAE-S
Logistic regression (pixel)
TabNet (pixel)
TempCNN
TempCNN, gated
XGBoost (field)
5-seed ensemble (100%)