


1 Graphical Abstract

2 Sparse Feature Selection, Not Deep Learning Capacity, Drives Spatial 3 Transfer in Crop Classification

4 Michael L. Mann, Sairam Venkatachalam, Disha Kacha, Devarsh Sheth, Ryan
5 Engstrom, Amir Jafari



The ranking inverts under spatial transfer

Dense temporal nets lead in-region, then collapse;
sparse feature-selection models rise to the top.

- Dense temporal net ▼ falls
- Linear baseline (robust)
- Sparse feature selection ▲ rises
- L-TAE-S (ours)

6 Sparse Feature Selection, Not Deep Learning Capacity, 7 Drives Spatial Transfer in Crop Classification

8 Michael L. Mann^{a,*}, Sairam Venkatachalam^b, Disha Kacha^b, Devarsh Sheth^b,
9 Ryan Engstrom^a, Amir Jafari^b

*^aDepartment of Geography & Environment, The George Washington
University, Washington, DC, 20052, USA*

^bData Science Program, The George Washington University, Washington, DC, 20052, USA

10 Abstract

Accurate crop maps underpin food-security monitoring and agricultural land-use planning across the Global South, yet operational mapping there means applying a model to regions where no labels were collected. Most studies instead validate within a single region (random or field-wise cross-validation), which measures in-distribution fit rather than spatial transfer. On a South African Sentinel-2 dataset (optical only), we benchmark twenty-three classical and deep models twice: in-region and on a spatial holdout tile. The rankings invert. In-region the dense temporal deep nets lead (macro-F1 0.72–0.78); under transfer this cluster fractures, as pixel-level L-TAE and a Transformer encoder drop from 0.78 to 0.58 and CNN-BiLSTM falls the most (0.72 to 0.46), while a TabNet ensemble becomes the best model (0.60) and simple trees and logistic regression (0.56) prove robust. The results suggest a mechanistic explanation: models with sparse, axis-aligned feature selection are more stable under transfer. TabNet inherits it through sparsemax masks, while dense nets discard it. A control isolates it: a Transformer with far greater attention capacity transfers no better than L-TAE, so capacity is not the lever; sparsity is. We graft that bias onto a deep network, L-TAE-S, which ties TabNet for the best spatial-transfer macro-F1, and its field-level variant has the smallest in-region-to-holdout gap of any model. Field-level aggregation is a second, substitutable lever, and robust models hold near-full accuracy on a quarter of the fields. For agencies mapping unlabeled terrain, the message is twofold: model selection in crop mapping is an artifact of the validation protocol, and the property deciding whether a model survives transfer is sparse, axis-aligned feature selection.

11 *Keywords:* Crop classification, deep learning, domain generalization, feature
12 selection, remote sensing, Sentinel-2, spatial transfer, time-series classification

13 1. Introduction

14 Accurate and timely crop classification is foundational to precision agri-
15 culture and to the agencies that monitor food security and agricultural land

*Corresponding author.
Preprint submitted to Remote Sensing Applications, Society and Environment July 14, 2026
Email address: mmann123@gmail.com (Michael L. Mann)

use, enabling yield forecasting, resource optimization, and food-security assessment (Bégué et al., 2020; Ienco et al., 2017). The Sentinel-2 constellation, with its multispectral richness and roughly five-day revisit, has made dense multi-temporal optical imagery widely available. Yet reliable crop classification remains hard throughout Africa and the Global South: spectral similarity between crops, irregular planting and phenological cycles, persistent cloud cover, and rain-fed dryland management all degrade accuracy (Bégué et al., 2020; Mahlayeye et al., 2024). Classical methods such as Random Forest (RF) and XGBoost have proven effective and efficient on engineered features (Mahlayeye et al., 2024; Mann et al., 2026; Mann and Warner, 2017; Mann et al., 2019), but their performance hinges on hand-crafted feature engineering (Maxwell et al., 2018; Belgiu and Drăguț, 2016). Deep learning—recurrent and convolutional temporal models—reduces reliance on hand-crafted predictors and captures phenological structure that summary statistics miss (Zhu et al., 2017; Rußwurm and Körner, 2018; Pelletier et al., 2019), but its generalizability in complex, low-label settings remains uncertain (Pankajakshan et al., 2025).

A second, and for our purposes critical, gap concerns *how generalization is measured*. Much of the crop-classification literature reports accuracy under random or field-wise cross-validation drawn from a single region (Belgiu and Drăguț, 2016; Garnot and Landrieu, 2020; Pelletier et al., 2019; Rußwurm and Körner, 2018; Ienco et al., 2017). Such protocols hold out individual pixels or fields but not *space*: training and test observations share the same atmosphere, soils, management calendars, and acquisition geometry, so the reported numbers reflect in-distribution fit rather than transfer to new terrain. Because operational maps are produced by applying a model to regions where no labels were collected, this distinction matters, yet studies that re-rank models under a spatially disjoint holdout remain rare (Pankajakshan et al., 2025; Jin et al., 2019; Rustowicz et al., 2019). This differs from the bulk of the transfer-learning literature, which adapts a model using target-domain data—labels, unlabeled imagery, or self-supervised pre-training (Ma et al., 2024; Wang et al., 2023; Mai et al., 2025). Our setting is the no-target-access case (domain generalization): models are trained on source tiles and applied to the holdout directly, with no adaptation, so what determines transfer is a property intrinsic to the model rather than a target-specific tuning step.

The Western Cape of South Africa concentrates the operational challenges—smaller fields, rain-fed dryland systems, heterogeneous planting rotations, and persistent winter cloud cover—while offering an unusually rich, high-quality labeled dataset (Radiant Earth Foundation, 2024). This makes it a realistic test bed: rich enough to train and compare modern architectures fairly, yet hard enough to expose the failure modes that matter in practice. We evaluate CNN-BiLSTM hybrids, TabNet ensembles, lightweight temporal attention (L-TAE) and full Transformer-encoder networks, temporal-convolution (TempCNN) networks, and patch-based 2D/3D CNNs, benchmarked against XGBoost, LightGBM, and RF. Critically, we evaluate every model twice: under field-wise cross-validation *within* the training region, and on a spatially disjoint holdout tile that no training pixel touches. We restrict inputs to multi-temporal

62 Sentinel-2 optical imagery alone, to isolate the spectral-time-series signal and
63 maximize applicability where only optical data are available. All code is public
64 (Mann, 2025a).

65 This study makes five contributions. First, a multi-architecture bench-
66 mark (twenty-three classical, deep, and patch-based models) showing that in-
67 region versus spatial-transfer evaluation produce *inverted* rankings on the same
68 data. Second, empirical evidence that models with sparse, axis-aligned feature-
69 selection, the property behind two decades of RF and gradient-boosted-tree
70 dominance, are more stable when moved to a holdout tile, and TabNet inherits
71 it via sparsemax masks. Third, a controlled architectural comparison isolating
72 sparse feature selection rather than attention capacity as the lever, including
73 L-TAE-S, a sparse-attention encoder we introduce that ties TabNet for the
74 best ST macro-F1 and achieves the smallest transfer gap in its field-level vari-
75 ant. Fourth, identification of field-level aggregation as a second, substitutable
76 robustness lever that recovers transfer for dense models but is redundant for
77 sparse ones. Fifth, a data-efficiency analysis under spatial transfer showing
78 the most transferable models are also the most data-efficient, holding near-full
79 holdout accuracy on a quarter of the training fields; as a supporting exper-
80 iment, a controlled comparison of `xr_fresh` time-series features against raw
81 reflectance under spatial transfer. This leads to our central question: across
82 machine-learning, deep-learning, and hybrid architectures, which most reliably
83 classifies crops from multi-temporal Sentinel-2 optical imagery *when evaluated*
84 *on a spatially disjoint region*, and what model properties determine whether
85 in-distribution accuracy survives the transfer?

86 1.1. Related Work

87 Early crop classification relied on classical learners (RF, SVM, logistic re-
88 gression, gradient-boosted trees) operating on engineered descriptors—spectral
89 indices, temporal summaries—valued for efficiency, interpretability, and robust-
90 ness on structured data (Maxwell et al., 2018; Belgiu and Drăguț, 2016; Breiman,
91 2001; Chen and Guestrin, 2016; Ke et al., 2017), though their performance de-
92 pends on feature engineering that can limit transferability across agroecological
93 regions (Wardlow et al., 2007; Mahlayeye et al., 2024). Deep models learn rep-
94 resentations directly from reflectance sequences: CNNs capture local spectral-
95 spatial pattern (Kussul et al., 2017), and recurrent and attention architectures
96 model seasonal phenology (Ienco et al., 2017; Rußwurm and Körner, 2018; Gar-
97 not and Landrieu, 2020; Rußwurm and Körner, 2020). Hybrid designs combine
98 convolutional feature extraction with temporal modeling (Fan et al., 2023; Tu-
99 fail et al., 2025; Bandar and Coşkunçay, 2024), or use deep backbones as feature
100 extractors feeding classical classifiers (Espejo-Garcia et al., 2020; Koklu et al.,
101 2022; Yang et al., 2020). A recurring assumption is that richer learned repre-
102 sentations generalize better than compact engineered ones; whether this holds
103 under genuine spatial transfer, rather than within-region validation, is less es-
104 tablished and is a central question here. Architectures that explicitly target
105 cross-scenario generalization (e.g., Cropformer (Wang et al., 2023)) remain the
106 exception and are rarely benchmarked against classical and deep baselines on a

Study area: Western Cape, South Africa




Figure 1: Study area in the Western Cape, South Africa. Two adjacent Sentinel-2 tiles (34S_19E_258N, 34S_19E_259N) form the training region; the immediately adjacent 34S_20E_259N tile (dashed) is held out for spatial-transfer evaluation. Fields are colored by crop type; the tiles are spatially disjoint, so no pixel or field is shared between training and holdout.

107 common spatially disjoint holdout. We fill this gap with a systematic comparison
108 under *both* protocols, and additionally evaluate `xr_fresh` as an automated
109 time-series feature-extraction framework in a crop-classification setting (Mann,
110 2025b).

111 **2. Data**

112 The study uses a labeled field dataset distributed through Radiant Earth’s
113 MLHub (Radiant Earth Foundation, 2024) and used as the basis of the AI4FoodSecurity
114 (South Africa) challenge (AI4EO, 2021), in an agriculturally productive part
115 of the Western Cape (Fig. 1). Two adjacent Sentinel-2 tiles (34S_19E_258N,

116 34S_19E_259N) form the training region (4,150 labeled fields over 622 km²; mean
 117 field 15.0 ha); a third, immediately adjacent tile (34S_20E_259N) is reserved as a
 118 holdout for spatial-transfer evaluation (2,417 fields over 282 km²). The tiles do
 119 not overlap, so no pixel or field is shared, and transfer must contend with differ-
 120 ent soils, geometries, and acquisition conditions. Because the holdout lies within
 121 the same agroecological zone and growing calendar, this is a relatively *local* spa-
 122 tial shift, and the gaps we report are best read as a conservative lower bound on
 123 degradation under broader transfer. Throughout, “in-region” denotes a field-
 124 wise train/validation/test split within the training region; “spatial-transfer”
 125 (ST) or “holdout” denotes predictions on the disjoint tile.

126 We classified five crop types: wheat, canola, lucerne/medics (alfalfa), barley,
 127 and small-grain grazing. The classes are heavily imbalanced—lucerne/medics
 128 is 43% of training fields and 56% of the holdout (Supplementary Material,
 129 Fig. S1(a))—which motivates macro-F1 as the primary metric (Sec. 3.3) and
 130 class-imbalance handling during training. Spectrally, wheat and barley overlap
 131 strongly, the dominant source of persistent confusion (Supplementary Mate-
 132 rial, Fig. S2(b)). Sentinel-2 Level-2A surface reflectance was acquired for the
 133 January–December 2017 growing season, cloud-masked with `s2cloudless` plus
 134 a shadow mask, and composited monthly. Six features per month are used—
 135 bands B2, B6, B11, B12, plus EVI and Hue, with EVI and Hue computed per
 136 scene before compositing—chosen to span the most informative spectral regions
 137 while limiting redundancy among correlated bands (Mann et al., 2026; Tufail
 138 et al., 2025). June 2017 yielded no usable imagery and May was discarded
 139 for cloud, leaving 10 monthly observations, so each pixel is a 60-dimensional
 140 time series (6 features $\times$ 10 months). Field polygons were digitized from sub-
 141 meter aerial photography and carry a single seasonal crop label (Radiant Earth
 142 Foundation, 2024); all predictions are made at the field level.

143 3. Methodology

144 We compare three paradigms: feature-based classical machine learning, representation-
 145 learning deep networks, and patch-based spatial models. To prevent leakage,
 146 all train/validation/test splitting is performed at the field level with a fixed
 147 random seed, so no field’s pixels straddle partitions; pixel- and patch-level pre-
 148 dictions are aggregated to a field label by majority vote (or, for probabilistic
 149 models, mean pooling). Computational cost across paradigms is reported in the
 150 Supplementary Material (Sec. S-C).

151 3.1. Classical Machine Learning

152 We automate feature extraction with the `xr.fresh` toolkit, computing a set
 153 of statistical and temporal descriptors (e.g., autocorrelation, linear trend, day-
 154 of-year of the seasonal maximum, quantiles, complexity) from each pixel’s 10-
 155 month profile, yielding a 174-dimensional per-pixel vector (Mann et al., 2026;
 156 Mann, 2025b); the full feature list is in the Supplementary Material (Sec. S-
 157 A, Table S1). Four baselines—logistic regression, RF (300 trees), LightGBM,

and XGBoost (James et al., 2013; Breiman, 2001; Ke et al., 2017; Chen and Guestrin, 2016; Friedman, 2001)—are trained on this representation, all of them tree- or margin-based with built-in feature selection or shrinkage; standard formulations are reproduced in the Supplementary Material (Sec. S-A). We also report soft-voting and stacked ensembles (RF/XGBoost/LightGBM, with logistic-regression meta-learner), an Optuna-tuned field-level XGBoost, and a SMOTETomek-resampled variant; the *Stacking* and *SMOTE Stacked* pair differs only in resampling, isolating the effect of class rebalancing. At the field level, pixel features are mean-pooled within each polygon before classification (Peña-Barragán et al., 2011).

3.2. Deep Learning Architectures

At the pixel level each observation is a 60-feature temporal sequence. Although the pixel models train on millions of pixels, pixels within a field are highly autocorrelated and share one label, so the number of statistically independent units is closer to the number of fields; field-wise splitting keeps evaluation honest at the field scale. We evaluate, as five-seed ensembles, a CNN-BiLSTM hybrid trained with weighted focal loss; TabNet (Arik and Pfister, 2021; Muruganathan et al., 2024), which treats the observations as tabular input and, at each decision step, uses an attentive transformer to produce a *sparse* feature mask via *sparsemax*, selecting a handful of features and ignoring the rest; the L-TAE (Garnot and Landrieu, 2020), which applies a single learned query over the temporal sequence; a standard multi-head Transformer encoder (Vaswani et al., 2017; Rußwurm and Körner, 2020); and TempCNN (Pelletier et al., 2019). Detailed equations for these standard models, and the field-level hybrid-voting rule, are in the Supplementary Material (Sec. S-A). The Transformer is included as an informative dense-attention comparison: it increases attention capacity relative to L-TAE while retaining no explicit sparse feature-selection mechanism. Patch-based models (a multi-channel 2D CNN and a 3D CNN) tile each field into 100 m patches resampled to 128×128 and are likewise trained as five-seed ensembles; details in the Supplementary Material (Sec. S-A).

3.2.1. Sparse-Attention L-TAE (L-TAE-S)

To test whether a tree-like, sparse inductive bias can be grafted onto a temporal-attention encoder, we introduce L-TAE-S, which inserts a TabNet-inspired channel-selection gate before the attention heads (Fig. 2). For each of the $n_{\text{head}} = 16$ heads, a per-head linear projection of the embedded, positionally-encoded sequence ($d_{\text{model}} = 128$, key dimension $d_k = 8$) is passed through *sparsemax* (Martins and Astudillo, 2016) to produce a sparse, instance-dependent mask over embedding channels. Formally, for head h and time step $\mathbf{z}_t \in \mathbb{R}^E$,

$$\mathbf{m}_{h,t} = \text{sparsemax}(\mathbf{P}_{h,t} \odot \text{BN}(W_h \mathbf{z}_t)), \quad \tilde{\mathbf{z}}_{h,t} = \mathbf{m}_{h,t} \odot \mathbf{z}_t,$$

$$\mathbf{P}_{h+1,t} = \mathbf{P}_{h,t} \odot \max(\gamma - \mathbf{m}_{h,t}, 0), \quad \mathbf{P}_{1,t} = \mathbf{1};$$

head h then forms its keys and values from $\tilde{\mathbf{z}}_{h,t}$ before applying L-TAE attention. Because *sparsemax* projects onto the probability simplex, most channels of $\mathbf{m}_{h,t}$

are exactly zero, and the relaxation prior $\mathbf{P}_{h,t}$ (coefficient $\gamma = 1.5$) discourages successive heads from re-selecting the same channels; a penalty $\lambda \sum_{h,t} \|\mathbf{m}_{h,t}\|_1$ ($\lambda = 10^{-3}$) adds mild additional pressure. This single block—the only architectural difference from the L-TAE—grafts a tree-like sparse feature-selection bias onto the attention backbone while keeping its temporal modeling, mirroring TabNet’s sparsemax masks (Arik and Pfister, 2021). We evaluate L-TAE-S at both the pixel and field level as a five-seed ensemble under the same weighted-focal objective and field-aggregation protocol as the other temporal networks.

3.3. Performance Metrics

Because the classes are severely imbalanced, plain accuracy is misleading. We adopt macro-F1 ($F1_m$, the unweighted mean of per-class F1) as the primary metric, since it weights every crop equally and penalizes failure on the minority cereals where models differ most. We report three further metrics: Cohen’s Kappa; weighted F1 (wF1, the support-weighted mean of per-class F1); and cross-entropy (Xent), the field-level scoring metric of the AI4FoodSecurity challenge (AI4EO, 2021) from which the dataset is drawn (lower is better, computed on hard label predictions). Formal definitions are in the Supplementary Material (Sec. S-B).

3.4. Quantifying Uncertainty

We quantify two sources of variability. For sampling uncertainty in the finite holdout, we apply a nonparametric bootstrap over the 2,417 holdout fields: for each model we resample its per-field predictions with replacement $B = 10,000$ times and recompute macro-F1, reporting the 2.5th–97.5th percentile interval as a 95% confidence interval (CI); For between-model gaps, we compute paired bootstrap differences by resampling holdout fields and subtracting model scores within each bootstrap replicate. For optimization noise, the per-seed standard deviation of ST macro-F1 across the five ensemble members is 0.015–0.037 for the headline temporal models, which is why we report five-seed ensembles throughout. Both sources are small relative to the ranking effects below: every field-bootstrap CI spans only $\approx \pm 0.02$ – 0.03 macro-F1, so we treat ST macro-F1 differences below ≈ 0.05 as not statistically distinguishable.

4. Results

We evaluate every model under two protocols: 1) in-region field-wise validation and 2) spatial transfer (ST) to the disjoint holdout tile—all reported metrics are field-level. Pixel- and patch-level models are aggregated to field labels, whereas field-level models are trained and evaluated on field-aggregated inputs




Figure 2: Architecture of the sparse-attention L-TAE-S. Before the multi-head temporal attention, a TabNet-style sparse feature gate (highlighted) applies a per-head **sparsemax** mask that selects a small subset of embedding channels, with a relaxation prior ($\gamma = 1.5$) discouraging heads from re-selecting the same channels. This block is the only architectural difference from the L-TAE; five seeds are ensembled and aggregated to field level.

4.1. In-Region Validation Overstates Performance and Misranks Models

Table 1 reports both protocols for all models, ranked by ST macro-F1, with the gap $\Delta = F1_m^{ST} - F1_m^{in-region}$. Two facts stand out. First, *every* model loses accuracy under transfer, but unevenly: from -0.04 (field-level L-TAE-S) to -0.26 (CNN-BiLSTM at the pixel level). Second, this unevenness reorders the models. In-region the dense temporal pixel nets cluster at the top (L-TAE 0.78, Transformer 0.78, L-TAE-S 0.77, TempCNN 0.77, CNN-BiLSTM 0.72); a practitioner choosing by in-region F1 would pick from this cluster. On the holdout the cluster fractures: the two in-region leaders (L-TAE pixel and the Transformer, $\Delta = -0.20$) drop to 0.58, TempCNN to 0.56, and CNN-BiLSTM ($\Delta = -0.26$) to 0.46, among the lowest absolute scores of any model.

The winners under transfer are a different group. The best single model is the TabNet pixel ensemble ($F1_m = 0.60$, $\kappa = 0.54$), which also posts the lowest holdout cross-entropy (4.42) of any model, tied on macro-F1 by both sparse-attention L-TAE-S variants (0.60). The most *robust* models, those with the smallest gaps, are the sparse learners: field-level L-TAE-S posts the single smallest gap ($\Delta = -0.04$), just ahead of logistic regression (-0.05) and the pixel-level tree learners. That L-TAE-S—a deep temporal attention network whose only twist is a learned gate restricting each head to a small channel subset—sits atop the robustness ranking is itself evidence for our thesis. The field bootstrap confirms which differences are real: the best transferers beat CNN-BiLSTM by a wide, significant margin (TabNet over CNN-BiLSTM pixel, $\Delta = 0.14$, 95% CI $[0.11, 0.18]$, $p < 0.001$), so the inversion is not holdout noise; at the very top the leaders (TabNet and both L-TAE-S variants) are statistically indistinguishable (all 0.60; TabNet vs. L-TAE-S, $p = 0.80$). We therefore do not claim a single best model; we claim that a sparse-feature-selection group occupies the top of the ST ranking while the dense temporal nets prized by in-region validation fall significantly below it.

Table 1: In-region versus spatial-transfer performance of all twenty-three models.

Model	Level	Features	In-reg		Spatial transfer (ST)				
			$F1_m$	Δ	$F1_m$	95% CI	κ	wF1	Xent
TabNet (raw)	pixel	raw	0.73	-0.13	0.60	.58-.63	0.54	0.70	4.42
L-TAE-S	field	raw	0.64	-0.04	0.60	.58-.62	0.52	0.69	4.77
L-TAE-S	pixel	raw	0.77	-0.17	0.60	.57-.62	0.53	0.70	4.63
L-TAE	field	raw	0.67	-0.07	0.59	.57-.62	0.47	0.67	5.51
L-TAE	pixel	raw	0.78	-0.20	0.58	.55-.60	0.49	0.68	5.20
Transformer	pixel	raw	0.78	-0.20	0.58	.55-.60	0.50	0.69	4.94
XGBoost	field	xr_fresh	0.70	-0.13	0.56	.54-.59	0.47	0.68	4.89
TempCNN	pixel	raw	0.77	-0.21	0.56	.54-.59	0.49	0.68	5.09
LightGBM	pixel	xr_fresh	0.62	-0.06	0.56	.54-.59	0.47	0.66	4.85
Logistic Regression	pixel	xr_fresh	0.61	-0.05	0.56	.54-.59	0.47	0.67	4.93
LightGBM	field	xr_fresh	0.68	-0.12	0.56	.53-.58	0.45	0.66	4.99
Voting	pixel	xr_fresh	0.69	-0.13	0.56	.53-.58	0.48	0.66	4.85
Random Forest	pixel	xr_fresh	0.65	-0.10	0.55	.52-.57	0.45	0.65	5.07
XGBoost	pixel	xr_fresh	0.64	-0.09	0.54	.52-.57	0.47	0.66	5.01
TempCNN	field	raw	0.67	-0.14	0.54	.51-.56	0.43	0.64	6.07
Stacking	field	xr_fresh	0.64	-0.11	0.53	.51-.56	0.48	0.67	5.32
CNN-BiLSTM	field	raw	0.64	-0.14	0.51	.48-.53	0.33	0.54	7.86

Table 2: Pixel- versus field-level spatial transfer.

Architecture	Pixel		Field		Field input
	F1 _m	Δ	F1 _m	Δ	
L-TAE-S	0.60	−0.17	0.60	− 0.04	raw-avg
L-TAE	0.58	−0.20	0.59	−0.07	raw-avg
XGBoost	0.54	−0.09	0.56	−0.13	xr_fresh
LightGBM	0.56	− 0.06	0.56	−0.12	xr_fresh
TempCNN	0.56	−0.21	0.54	−0.14	raw-avg
CNN-BiLSTM	0.46	−0.26	0.51	−0.14	raw-avg
TabNet	0.60	−0.13	0.43	−0.19	raw-avg

Architectures run at both granularities, ranked by field-level ST F1_m. The *Field input* column gives the field-level input (**xr_fresh** statistics for the trees, field-averaged raw sequence for the deep models). Field aggregation helps the dense temporal nets, is roughly neutral for the tree learners, and degrades TabNet. Bold marks the best value per column.

Model	Level	Features	F1 _m	Δ	F1 _m	95% CI	κ	wF1	Xent
3D CNN	patch	raw	0.70	−0.20	0.50	.48–.53	0.29	0.54	7.98
SMOTE Stacked	field	xr_fresh	0.69	−0.21	0.49	.46–.51	0.43	0.63	5.37
CNN-BiLSTM	pixel	raw	0.72	−0.26	0.46	.44–.48	0.32	0.56	7.73
Multi-Channel 2D CNN	patch	raw	0.63	−0.18	0.45	.42–.47	0.30	0.54	7.62
TabNet (raw, field-avg)	field	raw	0.62	−0.19	0.43	.41–.46	0.31	0.56	7.35
TabNet (xr_fresh)	field	xr_fresh	0.51	−0.15	0.37	.34–.40	0.25	0.53	6.13

Ranked by ST macro-F1. In-region uses field-wise validation; ST uses the disjoint holdout tile (34S_20E_259N). Δ is the gap; the 95% CI is the percentile bootstrap on ST macro-F1 from 10,000 resamples of the 2,417 holdout fields. Xent is the AI4FoodSecurity field-level scoring metric (lower is better). *Features* marks the automated **xr_fresh** statistics (174 features) versus raw reflectance. Bold marks the best value per column.

4.2. Pixel Versus Field Level: Aggregation as Variance Reduction

Several architectures were run at both the pixel and the field level (averaging per-pixel sequences within a field before classification). Table 2 shows that field-level aggregation acts as a variance-reduction regularizer for the *dense* temporal networks: L-TAE’s gap shrinks from −0.20 to −0.07, and CNN-BiLSTM, the worst pixel-level transferer, recovers most, its gap nearly halving from −0.26 to −0.14 (F1_m 0.46 → 0.51). The sparse-attention L-TAE-S follows the same dense-backbone pattern despite its channel gating: its gap shrinks from −0.17 to −0.04 while ST F1_m holds at 0.60, so the response to aggregation appears to depend on both the temporal backbone and the input granularity; the sparse gate does not eliminate the benefit of field-level aggregation. TabNet is the informative exception: aggregating its raw inputs to the field level *hurts* (F1_m 0.60 → 0.43), because its robustness already derives from sparse feature selection and field-level training starves it of the per-pixel observation volume it exploits. The tree learners transfer well at either level, and their pixel→field change is essentially neutral under transfer (XGBoost 0.54 → 0.56). In short, dense models need aggregation to generalize; sparse, tree-like models do not.




Figure 3: Per-class F1 on the holdout region for four representative models (TabNet, L-TAE-S, L-TAE, and CNN-BiLSTM). TabNet, L-TAE-S, and L-TAE classify lucerne/medics and canola well; the cereals (barley, wheat) and small-grain grazing concentrate the errors. CNN-BiLSTM’s accuracy drops on every class. Error bars are 95% bootstrap CIs over holdout fields.

286 4.3. Per-Class Performance and Persistent Confusions

287 Figure 3 breaks accuracy down crop by crop on the holdout for four represen-
 288 tative models: the best transferer (TabNet), our sparse-attention L-TAE-S, its
 289 dense counterpart L-TAE, and CNN-BiLSTM (the largest drop under transfer).
 290 The errors are highly uneven. All three robust models classify the dominant
 291 lucerne/medics and the distinctive canola well (TabNet F1 0.88 and 0.81; L-
 292 TAE-S 0.87 and 0.76; L-TAE 0.84 and 0.66), but every model struggles with the
 293 cereals and small-grain grazing: even TabNet reaches only 0.55 on barley, 0.48
 294 on wheat, and 0.30 on small-grain grazing. Adding the sparsity gate (L-TAE →
 295 L-TAE-S) on the same backbone improves the easier classes (canola 0.66 → 0.76)
 296 while leaving the hard cereals essentially unchanged (barley 0.50 → 0.50, wheat
 297 0.48 → 0.47). CNN-BiLSTM, by contrast, loses accuracy on the easy classes
 298 too (lucerne/medics 0.71, canola 0.65), showing its failure is a general inability
 299 to transfer, not a minority-class problem alone. The wheat–barley confusion is
 300 consistent with their spectral overlap and appears across every model.

301 4.4. Data Efficiency Under Spatial Transfer

302 The most transferable models are also the most data-efficient. We retrain
 303 a representative set on random 75%, 50%, and 25% subsets of the training




Figure 4: Data efficiency under spatial transfer: macro-F1 versus fraction of training fields retained, single training seed, each point the mean over five independent subsample draws (error bars ± 1 s.d.). The deep temporal models (L-TAE, L-TAE-S) are nearly flat to a quarter of the fields; the tree, linear, and TabNet baselines decline modestly. Per-fraction variation is within draw-to-draw sampling noise—largest for L-TAE-S (± 0.06 at 50%)—so single-split “dips” are not real. Absolute values run ~ 0.02 – 0.05 below the five-seed levels of Table 1 because a single training seed is used.

fields and re-evaluate on the holdout; because a single subsample can be unrepresentative, we repeat each fraction over five independent draws and plot the mean (Fig. 4). Averaged this way, the deep temporal models are essentially flat from full data down to a quarter of the fields (L-TAE and L-TAE-S vary by under 0.02 macro-F1 across fractions), while the tree, linear, and TabNet baselines lose only ~ 0.03 – 0.05 —modest given a four-fold cut in training data. The per-fraction wiggles seen in any single split are within draw-to-draw sampling variance: L-TAE-S, whose sparse channel gate makes it the most sensitive to *which* fields are sampled, swings ± 0.06 macro-F1 at the 50% fraction (versus ≤ 0.02 for the robust learners), so the apparent 50% “dip” of any one draw is sampling noise rather than a data-efficiency effect. Holdout accuracy is therefore limited far more by spatial transfer than by training-set size, reinforcing that what governs holdout accuracy is the model’s inductive bias rather than the volume of training data.

Table 3: Raw reflectance versus `xr_fresh` for a single fixed XGBoost learner.

Representation	Split	$F1_m$	wF1	κ
Raw reflectance	In-region	0.73	0.80	0.71
<code>xr_fresh</code>	In-region	0.72	0.76	0.68
Raw reflectance	Spatial transfer	0.59	0.69	0.50
<code>xr_fresh</code>	Spatial transfer	0.56	0.69	0.50

Identical FID-wise splits, preprocessing, and field-level mean pooling; only the input representation differs. Metrics are field-level macro-F1 ($F1_m$), weighted F1 (wF1), and Cohen’s κ .

4.5. Raw Reflectance versus `xr_fresh`

The best-transferring classical learners in Table 1 are exactly those fed `xr_fresh` statistics, so representation and model family are confounded there. To separate them, we trained a single XGBoost learner once on raw monthly reflectance and once on the per-pixel `xr_fresh` statistics, holding the FID-wise splits, preprocessing, and field-level mean pooling fixed (Table 3). The two representations are essentially indistinguishable under transfer ($F1_m = 0.59$ raw vs. 0.56 `xr_fresh`, a difference within the holdout bootstrap interval of Table 1; equal weighted-F1 and κ), and raw reflectance is, if anything, marginally ahead in-region as well. The `xr_fresh` representation therefore yields no measurable transfer-accuracy gain over raw reflectance for the trees; its value is operational—a compact, fast-to-fit, agronomically interpretable vector that trains in seconds (Supplementary Material, Sec. S-C)—rather than an accuracy edge. This is, to our knowledge, the first controlled comparison of `xr_fresh` against raw reflectance under spatial transfer.

5. Discussion

The conclusion about which model is “best” depends almost entirely on whether evaluation is in-region or under spatial transfer, and a model’s transfer behavior is predictable from its inductive bias.

5.1. Why Tree-Like Inductive Bias Transfers

The clearest pattern is that transfer is strongly associated with inductive bias, and not simply with model family or attention capacity. Figure 5 groups every model by inductive bias and plots its ST gap: linear models lose the least (mean $\Delta = -0.05$), tree-based models little (mean -0.11), the sparse-attention L-TAE-S sits essentially with the trees (mean -0.11 across its variants), while dense temporal/patch nets lose the most (mean -0.21) and SMOTE-augmented stacking is similarly fragile (-0.21). RF and gradient-boosted trees have been the remote-sensing workhorse for two decades not because they top in-region accuracy (dense nets often do) but because they survive when a model is applied to an unseen region (Belgiu and Drăguț, 2016; Maxwell et al., 2018)—the operational case. A tree decides one feature at a time, choosing the few features that

matter for each split and ignoring the rest; this sparsity may reduce reliance on
 region-specific quirks of phenology, soil, and atmosphere that pull richer mod-
 els off course. Independent corroboration spans settings: a 3D-CNN collapses
 under cross-sensor hyperspectral transfer while a sparse SVM holds (Pankajak-
 shan et al., 2025); RF matches or exceeds a dense CNN-LSTM on label-scarce
 Ghana fields (Rustowicz et al., 2019); a maize model drops 16 points moving
 across East Africa (Jin et al., 2019); and bagged/boosted trees out-transfer an
 elaborate instance-reweighting scheme from the U.S. to northwest China (Mai
 et al., 2025). TabNet, the best transferer here, is a neural network built to
 behave like a tree: a small learned gate picks a handful of features at each step
 (Arik and Pfister, 2021), so it clusters with the trees and inherits their transfer-
 ability. By contrast, CNN-BiLSTM, TempCNN, L-TAE, and the Transformer
 consume the whole temporal sequence at once with no feature selection. At
 the pixel level this leaves them fitting the training region tightly: they post the
 largest in-region-to-holdout gaps in the benchmark, with CNN-BiLSTM drop-
 ping furthest—consistent with evidence that the dates such models attend to
 are themselves contingent on the training region (Obadić et al., 2022). Dense
 attention is not intrinsically poor under spatial transfer: the pixel-level L-TAE
 loses 0.20 macro-F1, but the field-level L-TAE reduces this gap to 0.07 and the
 other temporal nets (Sec. 4.2), so it is the pixel-level dense models specifically
 that transfer worst.

The Transformer is a particularly clean control: although it deploys the
 most expressive attention of any model we test, far richer than L-TAE’s single
 query, at the pixel level it transfers no better than the pixel-level L-TAE (both
 $F1_m = 0.58$, $\Delta = -0.20$). Within the architectures evaluated, increasing atten-
 tion capacity alone did not improve transfer, whereas adding feature *selection*
 does: L-TAE-S and TabNet, the two models that explicitly select features before
 computing with them, are the best transferers. This isolates feature selection—
 not the attention mechanism, the temporal backbone, or model size—as the
 property governing transfer. The L-TAE-S experiment adds that this is *not* a
 tree-versus-deep tradeoff: the same selection step that drives tree splits can be
 grafted onto a dense architecture at the cost of a single block, and the payoff is
 large—L-TAE-S ties TabNet for the best holdout accuracy. That parity is not
 accuracy alone: at the pixel level L-TAE-S reaches TabNet’s best-in-class hold-
 out accuracy in roughly 40% of the per-seed training time (24 vs. 61 min), and
 its field-level variant is both faster and more accurate than any field-aggregated
 TabNet (Supplementary Material, Sec. S-C). The next generation of operational
 crop classifiers need not choose between the expressivity of deep learning and
 the transfer-robustness of trees.

5.2. Two Robustness Levers

Two caveats sharpen rather than weaken the story, and together they point
 to how we can choose the best models. First, it is *pixel-level* TabNet that wins;
 the field-aggregated TabNet variants are among the worst transferers (Table 1),
 because aggregation removes the per-pixel observation volume the sparse model
 exploits. Second, dense temporal attention transfers well *when* applied at the




Figure 5: Spatial-transfer gap (Δ macro-F1) grouped by inductive-bias family. Sparse/tree-like and linear models transfer (small $|\Delta|$); dense temporal/patch nets and synthetic augmentation overfit the training region. L-TAE-S, which adds sparse feature selection to a dense attention backbone, tracks the tree-based models. Bars are family means; points are individual models.

field level: L-TAE’s pixel-level gap shrinks substantially under aggregation, and L-TAE-S ends with an even smaller gap (Table 2). These point to two distinct, partly substitutable pathways. The first lever is built-in feature selection—the model itself picking a few features at a time (trees, TabNet, and now L-TAE-S); this is *intrinsic* robustness paid for at design time. The second is variance reduction imposed from outside via field-level aggregation, which averages away region-specific pixel noise before the classifier sees it—*external* robustness available even when the architecture is fixed. A model with the first lever gains nothing, and may lose observation volume, from the second; this is exactly why TabNet prefers pixels. These two properties can be combined. L-TAE-S adds a sparse gate to a temporal-attention backbone and then benefits from field aggregation, ending with the smallest transfer gap in the benchmark. Practitioners can read this as a flowchart: if free to choose or modify the model, build feature selection in (lever 1); if constrained to an off-the-shelf dense architecture, aggregate predictions to the field level (lever 2). Two further controlled comparisons confirm the ranking is not an artifact of ensembling or imbalance handling; both shift absolute scores but never reorder the robust and fragile groups, with the full tables in the Supplementary Material (Sec. S-D). Seed-ensembling gives only

a small, monotone lift that leaves the ordering intact—a single-seed L-TAE-S ($F1_m = 0.56$) already exceeds the five-seed CNN-BiLSTM and TabNet field ensembles (Table S4). And the imbalance-handling objective moves absolute scores but not the ranking: the robust L-TAE-S is almost invariant across five loss and sampler settings ($F1_m = 0.59$ – 0.60), whereas the fragile CNN-BiLSTM swings widely (0.41 – 0.56) and even at its best stays below L-TAE-S’s worst (Table S5).

5.3. Automated Features, Patch Models, and Persistent Difficulty

The classical pipelines paired with `xr_fresh` were not only robust but cheap, transferring as well as the best deep models (ST $F1_m$ 0.55 – 0.56) while training in minutes rather than hours (Supplementary Material, Sec. S-C); the controlled comparison in Sec. 4.5 confirms that this value is operational rather than an accuracy edge over raw reflectance. Patch-based convolutional models did not exceed pixel-level models and transferred poorly (3D CNN $F1_m = 0.50$; multi-channel 2D CNN 0.45): at 10 m resolution the native 100 m patches are only $\sim 10 \times 10$ pixels, so the spatial context is limited and resampling adds interpolation rather than detail, and a dense convolutional representation offers no transfer-protective sparsity. Finally, small-grain grazing, barley, and wheat are the hardest classes for every model (Fig. 3) and account for most accuracy lost under transfer; that even robust models struggle here, and that CNN-BiLSTM loses accuracy on every class, indicates the limitation is partly intrinsic spectral similarity and partly the transfer problem itself, not class imbalance alone.

5.4. Limitations

The study uses a single growing season (2017) and a single holdout tile that is spatially *adjacent* to the training tiles within the same agroecological zone. This is a relatively local shift, so the gaps we report most likely *understate* the degradation that fully cross-regional or cross-year transfer would produce; our numbers are best read as a conservative lower bound, and multi-region, multi-year holdouts are needed to confirm the inductive-bias ordering is stable under larger shifts. The labels derive from a public dataset rather than our own campaign, and we use Sentinel-2 optical data only—a deliberate scoping choice that nonetheless limits the achievable ceiling for optically similar cereals, making SAR fusion a clear next step.

6. Conclusion

In-region validation, the protocol the crop-classification literature reports almost universally, misranks twenty-three Sentinel-2 classifiers when they are applied to a spatially disjoint holdout. Dense temporal networks that lead under conventional evaluation fall sharply under transfer, the worst of them (CNN-BiLSTM) landing near the bottom, while parsimonious models such as tree ensembles, logistic regression, and a TabNet ensemble retain rank and accuracy. The mechanism is the same property that made tree ensembles the

remote-sensing default for two decades: a sparse, axis-aligned feature-selection bias that commits to a few robust signals and ignores the rest. TabNet inherits it through sparsemax masks; a Transformer with far greater attention capacity does not, and transfers no better than the lightweight L-TAE, pointing at sparsity, not capacity, as the lever. Grafting that lever onto an attention model (L-TAE-S) ties TabNet for the best transfer at a fraction of its training cost and, at the field level, posts the smallest in-region-to-holdout gap of any model. Field-level aggregation is a second, partly substitutable lever that helps with sparsity in L-TAE-S but hurts TabNet. We therefore recommend the following: evaluate on a spatially disjoint holdout before declaring one architecture better than another; prefer models with a built-in feature-selection bias when mapping unseen terrain; and, if dense temporal models are required, aggregate their predictions to the field level. The conditions that make the Western Cape difficult are the conditions under which agricultural remote sensing must operate across much of the Global South, and evaluating models the way they are actually deployed, applied to regions where no labels were collected, is what will make remote sensing deliver where it matters most.

CRedit authorship contribution statement

Michael L. Mann: Conceptualization, Methodology, Software, Formal analysis, Writing – original draft, Writing – review & editing, Supervision, Project administration. **Sairam Venkatachalam:** Software, Formal analysis, Writing – original draft, Investigation. **Disha Kacha:** Software, Formal analysis, Investigation, Writing – original draft. **Devarsh Sheth:** Software, Formal analysis, Investigation, Writing – original draft. **Ryan Engstrom:** Supervision, Methodology, Writing – review & editing. **Amir Jafari:** Methodology, Methodology, Supervision, Writing – review & editing.

Ethics statement

This article does not contain any studies with human participants or animals performed by any of the authors. The study is based solely on publicly available satellite imagery (Sentinel-2) and field-boundary crop labels, so no ethical approval was required.

Declaration of competing interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

488 Data availability

489 The processed dataset supporting this study is openly available on IEEE
490 DataPort at <https://dx.doi.org/10.21227/n69k-r458> (DOI: 10.21227/n69k-
491 r458). The source code is available at [https://github.com/mmann1123/South_](https://github.com/mmann1123/South_Africa_Crop_Comp)
492 [Africa_Crop_Comp](https://github.com/mmann1123/South_Africa_Crop_Comp) (Mann, 2025a). The study builds on the ground-reference
493 crop labels of the Western Cape Department of Agriculture and Radiant Earth
494 Foundation, “Crop Type Classification Dataset for Western Cape, South Africa”,
495 v1.0, Radiant MLHub (DOI: 10.34911/rdnt.j0co8q), used under the CC-BY-4.0
496 license; Sentinel-2 imagery was extracted via Google Earth Engine.

497 Acknowledgements

498 This research received no external funding. We thank the Western Cape
499 Department of Agriculture and the Radiant Earth Foundation for collecting
500 and openly distributing the labeled field dataset.

501 References

- 502 AI4EO, 2021. AI4FoodSecurity challenge (South Africa). European Space
503 Agency (ESA), in partnership with Planet, TUM, DLR and Radiant Earth,
504 <https://platform.ai4eo.eu/ai4food-security-south-africa>.
- 505 Arik, S.Ö., Pfister, T., 2021. TabNet: Attentive interpretable tabular learning.
506 Proceedings of the AAAI Conference on Artificial Intelligence 35, 6679–6687.
507 doi:10.1609/aaai.v35i8.16826.
- 508 Bandar, A., Coşkunçay, A., 2024. Crop classification with attention based
509 BI-LSTM and temporal convolution neural network combination for remote
510 sensing BreizhCrop time series data. *Yüzüncü Yıl Üniversitesi Fen Bilimleri*
511 *Enstitüsü Dergisi* 29, 173–188. doi:10.53433/yyufbed.1335866.
- 512 Bégué, A., Leroux, L., Soumaré, M., Faure, J.F., Diouf, A.A., Augusseau, X.,
513 Touré, L., Tonneau, J.P., 2020. Remote sensing products and services in sup-
514 port of agricultural public policies in Africa: overview and challenges. *Front-*
515 *iers in Sustainable Food Systems* 4, 58. doi:10.3389/fsufs.2020.00058.
- 516 Belgiu, M., Drăguţ, L., 2016. Random forest in remote sensing: A review of
517 applications and future directions. *ISPRS Journal of Photogrammetry and*
518 *Remote Sensing* 114, 24–31. doi:10.1016/j.isprsjprs.2016.01.011.
- 519 Breiman, L., 2001. Random forests. *Machine Learning* 45, 5–32. doi:10.1023/A:
520 1010933404324.
- 521 Chen, T., Guestrin, C., 2016. XGBoost: A scalable tree boosting system,
522 in: Proceedings of the 22nd ACM SIGKDD International Conference on
523 Knowledge Discovery and Data Mining, pp. 785–794. doi:10.1145/2939672.
524 2939785.

Espejo-Garcia, B., Mylonas, N., Athanasakos, L., Fountas, S., Vasilakoglou, I.,
2020. Towards weeds identification assistance through transfer learning. *Computers and Electronics in Agriculture* 171, 105306. doi:10.1016/j.compag.2020.105306.

Fan, X., Chen, L., Xu, X., Yan, C., Fan, J., Li, X., 2023. Land cover classification of remote sensing images based on hierarchical convolutional recurrent neural network. *Forests* 14, 1881. doi:10.3390/f14091881.

Friedman, J.H., 2001. Greedy function approximation: A gradient boosting machine. *Annals of Statistics* 29, 1189–1232. doi:10.1214/aos/1013203451.

Garnot, V.S.F., Landrieu, L., 2020. Lightweight temporal self-attention for classifying satellite images time series, in: *Advanced Analytics and Learning on Temporal Data (AALTD 2020)*, ECML PKDD Workshops, Springer. pp. 171–181. doi:10.1007/978-3-030-65742-0_12.

Ienco, D., Gaetano, R., Dupaquier, C., Maurel, P., 2017. Land cover classification via multitemporal spatial data by recurrent neural networks. *arXiv preprint arXiv:1704.04055*.

James, G., Witten, D., Hastie, T., Tibshirani, R., 2013. *An Introduction to Statistical Learning*. Springer.

Jin, Z., Azzari, G., You, C., Di Tommaso, S., Aston, S., Burke, M., Lobell, D.B., 2019. Smallholder maize area and yield mapping at national scales with Google Earth Engine. *Remote Sensing of Environment* 228, 115–128. doi:10.1016/j.rse.2019.04.016.

Ke, G., Meng, Q., Finley, T., Wang, T., Chen, W., Ma, W., Ye, Q., Liu, T.Y., 2017. LightGBM: A highly efficient gradient boosting decision tree, in: *Advances in Neural Information Processing Systems (NeurIPS)*.

Koklu, M., Unlarsen, M.F., Ozkan, I.A., Aslan, M.F., Sabanci, K., 2022. A CNN-SVM study based on selected deep features for grapevine leaves classification. *Measurement* 188, 110425. doi:10.1016/j.measurement.2021.110425.

Kussul, N., Lavreniuk, M., Skakun, S., Shelestov, A., 2017. Deep learning classification of land cover and crop types using remote sensing data. *IEEE Geoscience and Remote Sensing Letters* 14, 778–782. doi:10.1109/LGRS.2017.2681128.

Ma, Y., Chen, S., Ermon, S., Lobell, D.B., 2024. Transfer learning in environmental remote sensing. *Remote Sensing of Environment* 301, 113924. doi:10.1016/j.rse.2023.113924.

Mahlayeye, M., Darvishzadeh, R., Nelson, A., 2024. Characterising maize and intercropped maize spectral signatures for cropping pattern classification. *International Journal of Applied Earth Observation and Geoinformation* 128, 103699. doi:10.1016/j.jag.2024.103699.

565 Mai, J., Feng, Q., Fu, S., Wang, R., Zhang, S., Zhang, R., Liang, T.,
566 2025. Enhancing crop type mapping in data-scarce regions through trans-
567 fer learning: A case study of the Hexi Corridor. *Remote Sensing* 17, 1494.
568 doi:10.3390/rs17091494.

569 Mann, M.L., 2025a. South_Africa_Crop_Comp: Source code for multitempo-
570 ral crop classification. GitHub repository, [https://github.com/mmann1123/](https://github.com/mmann1123/South_Africa_Crop_Comp)
571 [South_Africa_Crop_Comp](https://github.com/mmann1123/South_Africa_Crop_Comp).

572 Mann, M.L., 2025b. xr_fresh: Automated time series feature extraction for
573 remote sensing and gridded data. *Journal of Open Source Software* 10, 9009.
574 doi:10.21105/joss.09009.

575 Mann, M.L., Colson, L., Nealon, R., Engstrom, R., Nakacwa, S., 2026. Lite
576 learning: Efficient crop classification in Tanzania using feature extraction
577 with machine learning and crowd sourcing. *IEEE Journal of Selected Topics*
578 *in Applied Earth Observations and Remote Sensing* In press.

579 Mann, M.L., Warner, J.M., 2017. Ethiopian wheat yield and yield gap estima-
580 tion: A spatially explicit small area integrated data approach. *Field Crops*
581 *Research* 201, 60–74. doi:10.1016/j.fcr.2016.10.014.

582 Mann, M.L., Warner, J.M., Malik, A.S., 2019. Predicting high-magnitude, low-
583 frequency crop losses using machine learning: an application to cereal crops in
584 Ethiopia. *Climatic Change* 154, 211–227. doi:10.1007/s10584-019-02432-7.

585 Martins, A.F.T., Astudillo, R.F., 2016. From softmax to sparsemax: A sparse
586 model of attention and multi-label classification, in: *Proceedings of the 33rd*
587 *International Conference on Machine Learning (ICML)*, pp. 1614–1623.

588 Maxwell, A.E., Warner, T.A., Fang, F., 2018. Implementation of machine-
589 learning classification in remote sensing: an applied review. *International*
590 *Journal of Remote Sensing* 39, 2784–2817. doi:10.1080/01431161.2018.
591 1433343.

592 Muruganantham, P., Wibowo, S., Grandhi, S., Islam, N., 2024. A deep learning
593 approach for dealing with tabular data in crop classification, in: *Proceedings*
594 *of the IEEE International Conference on ICT Convergence (ICTC)*, pp. 2054–
595 2059. doi:10.1109/ICTC62082.2024.10826760.

596 Obadic, I., Roscher, R., Oliveira, D.A.B., Zhu, X.X., 2022. Exploring
597 self-attention for crop-type classification explainability. *arXiv preprint*
598 *arXiv:2210.13167*.

599 Pankajakshan, P., Padmasanan, A., Sundar, S., 2025. Deep architectures fail
600 to generalize: A lightweight alternative for agricultural domain transfer in
601 hyperspectral images. *Sensors* 26, 174. doi:10.3390/s26010174.

602 Pelletier, C., Webb, G.I., Petitjean, F., 2019. Temporal convolutional neural
603 network for the classification of satellite image time series. *Remote Sensing*
604 11, 523. doi:10.3390/rs11050523.

605 Peña-Barragán, J.M., Ngugi, M.K., Plant, R.E., Six, J., 2011. Object-based
606 crop identification using multiple vegetation indices, textural features and
607 crop phenology. *Remote Sensing of Environment* 115, 1301–1316. doi:10.
608 1016/j.rse.2011.01.009.

609 Radiant Earth Foundation, 2024. South Africa crop type competition. Source
610 Cooperative repository. Data DOI: 10.34911/rdnt.j0co8q.

611 Rußwurm, M., Körner, M., 2018. Multi-temporal land cover classification
612 with sequential recurrent encoders. *ISPRS International Journal of Geo-*
613 *Information* 7, 129. doi:10.3390/ijgi7040129.

614 Rußwurm, M., Körner, M., 2020. Self-attention for raw optical satellite time
615 series classification. *ISPRS Journal of Photogrammetry and Remote Sensing*
616 169, 421–435. doi:10.1016/j.isprsjprs.2020.06.006.

617 Rustowicz, R.M., Cheong, R., Wang, L., Ermon, S., Burke, M., Lobell, D., 2019.
618 Semantic segmentation of crop type in Africa: A novel dataset and analysis
619 of deep learning methods, in: *Proceedings of the IEEE/CVF Conference on*
620 *Computer Vision and Pattern Recognition (CVPR) Workshops*, pp. 75–82.

621 Tufail, R., Tassinari, P., Torreggiani, D., 2025. Deep learning applications for
622 crop mapping using multi-temporal Sentinel-2 data and red-edge vegetation
623 indices: Integrating convolutional and recurrent neural networks. *Remote*
624 *Sensing* 17, 3207. doi:10.3390/rs17183207.

625 Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A.N.,
626 Kaiser, L., Polosukhin, I., 2017. Attention is all you need, in: *Advances in*
627 *Neural Information Processing Systems (NeurIPS)*, pp. 5998–6008.

628 Wang, H., Chang, W., Yao, Y., Yao, Z., Zhao, Y., Li, S., Liu, Z., Zhang, X.,
629 2023. Cropformer: A new generalized deep learning classification approach
630 for multi-scenario crop classification. *Frontiers in Plant Science* 14, 1130659.
631 doi:10.3389/fpls.2023.1130659.

632 Wardlow, B.D., Egbert, S.L., Kastens, J.H., 2007. Analysis of time-series
633 MODIS 250 m vegetation index data for crop classification in the U.S. Central
634 Great Plains. *Remote Sensing of Environment* 108, 290–310. doi:10.1016/
635 j.rse.2006.11.021.

636 Yang, S., Gu, L., Li, X., Jiang, T., Ren, R., 2020. Crop classification method
637 based on optimal feature selection and hybrid CNN-RF networks for multi-
638 temporal remote sensing imagery. *Remote Sensing* 12, 3119. doi:10.3390/
639 rs12193119.

640 Zhu, X.X., Tuia, D., Mou, L., Xia, G.S., Zhang, L., Xu, F., Fraundorfer, F.,
641 2017. Deep learning in remote sensing: A comprehensive review and list of
642 resources. *IEEE Geoscience and Remote Sensing Magazine* 5, 8–36. doi:10.
643 1109/MGRS.2017.2762307.